



Profiling and Optimization

Lochan Khedkar

Outline



- What is Profiling
- Sampling based profiling
- Instrumentation based profiling
- Flow of Profiling and Analysis
- Profilers available for performance analysis
- OpenMP Profiling using Vtune
- Analysis of OpenMP application
- Optimization
- GNU Profiler Next Generation
- OpenMP analysis using GProfNG
- Tuning and Analysis Utility Profiler
- Profiling MPI application using Tuning and Analysis utility (TAU)Profiler



What is Profiling



- Application profiling is the process of analysing and measuring the performance of a software application in order to identify and diagnose performance issues or bottlenecks.
- Profiling tools help developers to identify areas of code that are taking a long time to execute, or that are using an excessive amount of system resources like memory, CPU or disk I/O.
- The purpose of profiling is to identify any bottlenecks or areas of the application that can be optimized to improve its performance, scalability, and reliability.
- By profiling their applications, developers can ensure that they are delivering high-quality software that performs well and meets the needs of their users.
- Profiling involves collecting data on various aspects of the application and its analysis.
- Implemented through either
 - Sampling based profiling
 - Instrumentation based Profiling



Sampling Based Profiling




- With sampling, the executable is stopped at regular intervals. Each time it is halted, key information is gathered and stored.
- At specified intervals, the Sampling method collects information about the functions that are executing in your application. While the program is interrupted the profiler grabs a snapshot of its current state.
- Sampling is a low cost to the profiler and has little effect on the execution of the application being profiled.
- If you need accurate measurements of call times or are looking for performance issues in an application, then sampling based profiler are useful.
- Sampling has less accuracy in the number of calls



Inclusive Vs Exclusive CPU Time

- This is an important concept in profiling tools
- The inclusive metric includes all callees underneath the caller - For example, all the CPU time accumulated when executing a function
- The exclusive metric excludes everything outside the caller - For example, the CPU time accumulated outside of calling other functions



Function	Inclusive time	Exclusive time
A	75	10
B	20	20
C	30	5
D	15	15
E	25	25



```
Void Alpha()  
{
```



30 samples

```
Beta();
```

```
}
```

```
Void Beta()  
{
```



50 samples

```
}
```


Functions	Inclusive	Exclusive
Alpha	80	30
Beta	50	50



Instrumentation based Profiling



- The first and earliest type are instrumenting profilers.
- Data collection is done by tools that either injecting code into a binary file that captures timing information or by using callback hooks to collect and emit exact timing and call count information while an application runs.
- The instrumentation method has a high overhead when compared to sampling-based approaches.
- Instrumentation profiling is that you can get exact call counts on how many times your functions were called.
- Instrumentation can be achieved during
 - Compile time Instrumentation
 - Binary Instrumentation



```
T=E*F;  
For (I=1;I<N;I++)  
{  
V[I]=C[I]*B[I];  
A[I]=C(2I+4);  
}
```

```
T=E*F;
```

Instrumentation code

```
For (I=1;I<N;I++)
```

```
{
```

```
V[I]=C[I]*B[I];
```


```
A[I]=C(2I+4);
```

```
}
```

Instrumentation code



Flow of Profiling and Analysis





Profilers Available for performance analysis



- Intel's Vtune
- GNU Profiler(gprof)
- GNU Profiler Next Generation(GProfNG) by Oracle
- Tuning and analysis Utility Profiler by Oregon University



OpenMP Profiling using Vtune

Installation steps for vtune

- `wget https://registrationcenter-download.intel.com/akdlm/IRC_NAS/56d0db2b-1ff1-4abe-857a-72ca9be22bd3/l_oneapi_vtune_p_2024.0.1.14_offline.sh`
- `chmod u+x l_oneapi_vtune_p_2024.0.1.14_offline.sh`
- `./l_oneapi_vtune_p_2024.0.1.14_offline.sh`



Hotspot analysis using Vtune



```
gcc -g -fopenmp -o mat_mult mat_mult.c
vtune -collect hotspots ./mat_mult
```

```
vtune: Executing actions 75 % Generating a report Elapsed Time: 54.413s
CPU Time: 31.760s
Effective Time: 31.110s
Spin Time: 0.640s
  Imbalance or Serial Spinning: 0.640s
  Lock Contention: 0s
  Other: 0s
Overhead Time: 0.010s
  Creation: 0.010s
  Scheduling: 0s
  Reduction: 0s
  Atomics: 0s
  Other: 0s
Total Thread Count: 20
Paused Time: 0s

Top Hotspots
Function      Module      CPU Time  % of CPU Time(%)
-----
printf        libc.so.6   25.906s   81.6%
multiply_omp_fn.0  mm         3.520s   11.1%
putchar       libc.so.6   1.644s    5.2%
gomp_simple_barrier_wait  libgomp.so.1  0.290s    0.9%
gomp_team_barrier_wait_end  libgomp.so.1  0.220s    0.7%
[Others]      N/A         0.180s    0.6%
Effective CPU Utilization: 0.4%
| The metric value is low, which may signal a poor logical CPU cores
| utilization caused by load imbalance, threading runtime overhead, contended
| synchronization, or thread/process underutilization. Explore sub-metrics to
| estimate the efficiency of MPI and OpenMP parallelism or run the Locks and
| Waits analysis to identify parallel bottlenecks for other parallel runtimes.
|
Average Effective CPU Utilization: 0.169 out of 40
Collection and Platform Info
Application Command Line: ./mm
Operating System: 3.10.0-1160.el7.x86_64 #@@#####@@@ #
```



Analysis



- Effective CPU execution time is 32.249 before optimization.
- Presence of spin time in Imbalance or serial spinning category indicates potential area for optimization.
- Also overhead time is crucial for optimal performance and should be as low as possible.
- Thread count should be appropriate for your workload and does not lead to excessive overhead.
- Logical core utilization value indicate poor CPU utilization.
- Low CPU utilization caused due to load imbalance, threading runtime overhead, synchronization and thread under utilization
- `printf` and `multiply._omp_fn.o` consume maximum CPU time and needs to be optimize.
- It indicates that parallel region `multiply._omp_fn.o` can be optimized further.



Optimization



- Use thread affinity to bind threads to specific CPU and avoid using large no of threads for reducing thread overheads.
- Experiment with chunk size in schedule clause to find the optimal balancing between load balancing and overheads.
- Consider using collapse clause to combine multiple nested loops into single loop for improved parallelization.
- Enable compiler optimization flags to allow compiler to perform additional optimization.
- Utilize SIMD (single instructions multiple data) by enabling compiler vectorization flags
- Experiment with loop unrolling to reduce loop overhead and improve instruction level parallelism.



Hotspot analysis using Vtune after optimization

```
vtune: Executing actions 75 % Generating a report Elapsed Time: 51.335s
CPU Time: 34.408s
Effective Time: 34.078s
Spin Time: 0.330s
  Imbalance or Serial Spinning: 0.330s
  Lock Contention: 0s
  Other: 0s
Overhead Time: 0s
  Creation: 0s
  Scheduling: 0s
  Reduction: 0s
  Atomics: 0s
  Other: 0s
Total Thread Count: 20
Paused Time: 0s

Top Hotspots
Function      Module      CPU Time  % of CPU Time(%)
-----
printf        libc.so.6   26.972s   78.4%
multiply_omp_fn.0 mm          5.388s   15.7%
putchar       libc.so.6   1.688s    4.9%
gomp_simple_barrier_wait libgomp.so.1 0.130s    0.4%
gomp_team_barrier_wait_end libgomp.so.1 0.100s    0.3%
[Others]      N/A         0.130s    0.4%
Effective CPU Utilization: 0.5%
| The metric value is low, which may signal a poor logical CPU cores
| utilization caused by load imbalance, threading runtime overhead, contended
| synchronization, or thread/process underutilization. Explore sub-metrics to
| estimate the efficiency of MPI and OpenMP parallelism or run the Locks and
| Waits analysis to identify parallel bottlenecks for other parallel runtimes.
|
Average Effective CPU Utilization: 0.214 out of 40
Collection and Platform Info
Application Command Line: ./mm
Operating System: 3.10.0-1160.el7.x86_64 #@@#####@#@# #
```



GNU Profiler Next Generation(GProfNG) by Oracle



Installation

- wget <https://ftp.gnu.org/gnu/binutils/binutils-2.42.tar.gz>
- tar -xvf [binutils-2.42.tar.gz](https://ftp.gnu.org/gnu/binutils/binutils-2.42.tar.gz)
- ./configure --prefix=/install_directory --with-gmp=/gmp_install_directory
- make ;make install



Openmp Analysis with GProfNG



- `gcc -g -fopenmp -o mat_mult mat_mult.c`
- Collect the data using following command
- `gprofng collect app ./mat_mult`
- Display the hotspot functions in `mat_mult.c` application using following command
- `gprofng display text -functions test.2.er`

```
[lochannk@login02 oneapi]$ gprofng display text -functions test.2.er/  
Functions sorted by metric: Exclusive Total CPU Time
```

Excl. Total CPU sec.	%	Incl. Total CPU sec.	%	Name
7.035	100.00	7.035	100.00	<Total>
3.342	47.51	3.342	47.51	multiply._omp_fn.0
2.822	40.11	2.822	40.11	<static>@0xefa99 (<libc-2.17.so>)
0.250	3.56	3.252	46.23	__printf_fp_l
0.120	1.71	3.422	48.65	vfprintf
0.100	1.42	0.100	1.42	__mpn_divrem
0.070	1.00	2.892	41.11	_IO_new_file_overflow
0.070	1.00	0.170	2.42	hack_digit.13666
0.060	0.85	0.060	0.85	__GI_memcpy
0.060	0.85	3.653	51.92	multiply
0.040	0.57	0.040	0.57	_IO_new_file_xsputn
0.040	0.57	0.040	0.57	gomp_team_barrier_wait_end
0.020	0.28	0.020	0.28	__overflow
0.020	0.28	0.020	0.28	gomp_barrier_wait_end
0.010	0.14	0.010	0.14	__GI_strlen
0.010	0.14	0.010	0.14	strchrnul
0.	0.	2.822	40.11	_IO_new_do_write
0.	0.	2.822	40.11	_IO_new_file_write
0.	0.	3.653	51.92	__libc_start_main
0.	0.	0.380	5.41	collector_root
0.	0.	0.020	0.28	gomp_team_end
0.	0.	0.380	5.41	gomp_thread_start
0.	0.	3.653	51.92	main
0.	0.	3.422	48.65	printf
0.	0.	0.150	2.13	putchar



Openmp thread analysis using GProfNG

- Threads analysis is done by following command
- gprofng display text –threads test.2.er/

```
[lochannk@login02 oneapi]$ gprofng display text -threads test.2.er/  
Objects sorted by metric: Exclusive Total CPU Time
```

Excl. Total CPU sec.	%	Name
7.035	100.00	<Total>
3.653	51.92	Process 1, Thread 1
3.342	47.51	Process 1, Thread 3
0.020	0.28	Process 1, Thread 2
0.020	0.28	Process 1, Thread 4

- gprofng display text –thread_select 3 –function

```
[lochannk@login02 oneapi]$ gprofng display text -thread_select 2 -functions test.2.er/  
Exp Sel Total  
=== === =====  
1 2 4  
Functions sorted by metric: Exclusive Total CPU Time
```

Excl. Total CPU sec.	%	Incl. Total CPU sec.	%	Name
0.020	100.00	0.020	100.00	<Total>
0.010	50.00	0.010	50.00	gomp_barrier_wait_end
0.010	50.00	0.010	50.00	gomp_team_barrier_wait_end
0.	0.	0.020	100.00	collector_root
0.	0.	0.020	100.00	gomp_thread_start



Profiling MPI application using Tuning and Analysis utility (TAU)Profiler



Installation

- `wget http://tau.uoregon.edu/tau.tgz`
- `tar -xvf tau.tgz`
- `/configure -c++=mpicxx -cc=mpicc -fortran=mpif90 -tag=openmpi -mpi -bfd=download -pdt=/home/lochannk/pdtoolkit-3.25.1 -papi=/opt/cray/pe/papi/6.0.0.2`
- `make`
- `make install`



MPI application functions analysis



- `mpicc -g -o matrix_mult mpi_mm.c`
- `mpirun -np 2 tau_exec -T openmpi ./matrix_mult`

```
[lochan@login02 oneapi]$ pprof
Reading Profile files in profile.*

NODE 0;CONTEXT 0;THREAD 0:
-----
%Time    Exclusive    Inclusive    #Call    #Subrs    Inclusive Name
      msec      total msec
-----
100.0      0.813        388          1          1    388802 .TAU application
 99.8         21        387          1         11    387989 taupreload_main
 60.1        233        233          1          0    233500 MPI_Init()
 24.5         95         95          1          0    95352 MPI_Finalize()
  7.7         30         30          3          0    10010 MPI_Recv()
  2.1          8          8          4          0    2006 MPI_Send()
  0.0      0.008      0.008          1          0         8 MPI_Comm_size()
  0.0      0.005      0.005          1          0         5 MPI_Comm_rank()

NODE 1;CONTEXT 0;THREAD 0:
-----
%Time    Exclusive    Inclusive    #Call    #Subrs    Inclusive Name
      msec      total msec
-----
100.0      0.625        386          1          1    386785 .TAU application
 99.8         30        386          1         11    386160 taupreload_main
 60.5        233        233          1          0    233876 MPI_Init()
 28.0        108        108          1          0    108387 MPI_Finalize()
  3.6         13         13          4          0     3442 MPI_Recv()
  0.0      0.102      0.102          3          0         34 MPI_Send()
  0.0      0.004      0.004          1          0         4 MPI_Comm_rank()
  0.0      0.003      0.003          1          0         3 MPI_Comm_size()
```



MPI performance counter analysis



- For performance counter analysis set export TAU_METRICS=" TIME,PAPI_TOT_CYC"
- `mpirun -np 2 tau_exec -T openmpi ./matrix_m`

```
[lochannk@login02 oneapi]$ cd MULTI__PAPI_TOT_CYC/
[lochannk@login02 MULTI__PAPI_TOT_CYC]$ pprof
Reading Profile files in profile.*

NODE 0;CONTEXT 0;THREAD 0:
-----
%Time   Exclusive   Inclusive   #Call   #Subrs Count/Call Name
      counts total counts
-----
100.0   3.835E+05   2.122E+08      1        1 212226018 .TAU application
 99.8   2.258E+07   2.118E+08      1       11 211842488 taupreload_main
 35.2   7.474E+07   7.474E+07      3        0 24914871 MPI_Recv()
 33.4   7.079E+07   7.079E+07      1        0 70791802 MPI_Init()
 13.1   2.788E+07   2.788E+07      4        0 6969219 MPI_Send()
  7.5   1.583E+07   1.583E+07      1        0 15830377 MPI_Finalize()
  0.0   1.117E+04   1.117E+04      1        0 11170 MPI_Comm_rank()
  0.0         5965         5965      1        0 5965 MPI_Comm_size()

NODE 1;CONTEXT 0;THREAD 0:
-----
%Time   Exclusive   Inclusive   #Call   #Subrs Count/Call Name
      counts total counts
-----
100.0   3.638E+05   1.896E+08      1        1 189624364 .TAU application
 99.8   7.45E+07   1.893E+08      1       11 189260527 taupreload_main
 37.3   7.077E+07   7.077E+07      1        0 70774385 MPI_Init()
 14.6   2.776E+07   2.776E+07      4        0 6941070 MPI_Recv()
  8.5   1.607E+07   1.607E+07      1        0 16065442 MPI_Finalize()
  0.1   1.383E+05   1.383E+05      3        0 46086 MPI_Send()
  0.0   1.149E+04   1.149E+04      1        0 11489 MPI_Comm_rank()
  0.0         6370         6370      1        0 6370 MPI_Comm_size()
```



THANK YOU !