

# HPC Architecture


Shamjith K V, Lochan Kedkar, Abhishek Patil

# Overview

- HPC Job Scheduling,
- Performance analysis
- Debugging
- Profiling,
- Petascale and Exascale Computing

# HPC Job Scheduling

# HPC Cluster



# Cluster Terminology

- Supercomputer/High Performance Computing (HPC) cluster: A collection of similar computers connected by a high speed interconnect that can act in concert with each other.
- Server, Node, Blade, Box, Machine : An individual motherboard with CPU, memory, network and local hard drive.
- CPU (Socket): Central Processing Unit, a single silicon die that can contain multiple computational cores
- Core: Basic unit of compute that runs a single instruction of code
- GPGPU: General Purpose Graphics Processing Unit, a GPU designed for supercomputing.
- InfiniBand (IB): A near zero latency high bandwidth interconnect used in Supercomputing
- Serial: Doing tasks/instructions in sequence on a single core
- Parallel: Doing tasks/instructions on multiple cores simultaneously
- I/O: Input/Output, a general term for reading and writing files to/from storage whether local or remote.

# What is a Job

- Job
  - user's program/name of an executable.
  - input data and parameters
  - environment variables
  - required libraries
  - descriptions of computing resources required
- Job Script
  - Formal specification
  - identifies an application to run along with its input data and environment variables
  - requests computing resources

# Local Resource Manager(LRM)/Batch System

- Job Scheduler or Workload Manager
  - Identifies jobs to run, selects the resources for the job, and decides when to run the job
- Resource Manager
  - identifies the compute resources and keeps track of their usage and feeds back this information to the workload manager
- Execution Manager
  - job initiation and start of execution is co-ordinated by the execution manager of the batch system.

# Job Scheduling


- The LRM is responsible for receiving and parsing the job script.
- If a job cannot be executed immediately, it is added to a queue.
- The job waits in the queue until the job's requested resources are available.
- The LRM then runs the job, when the requested resources are available



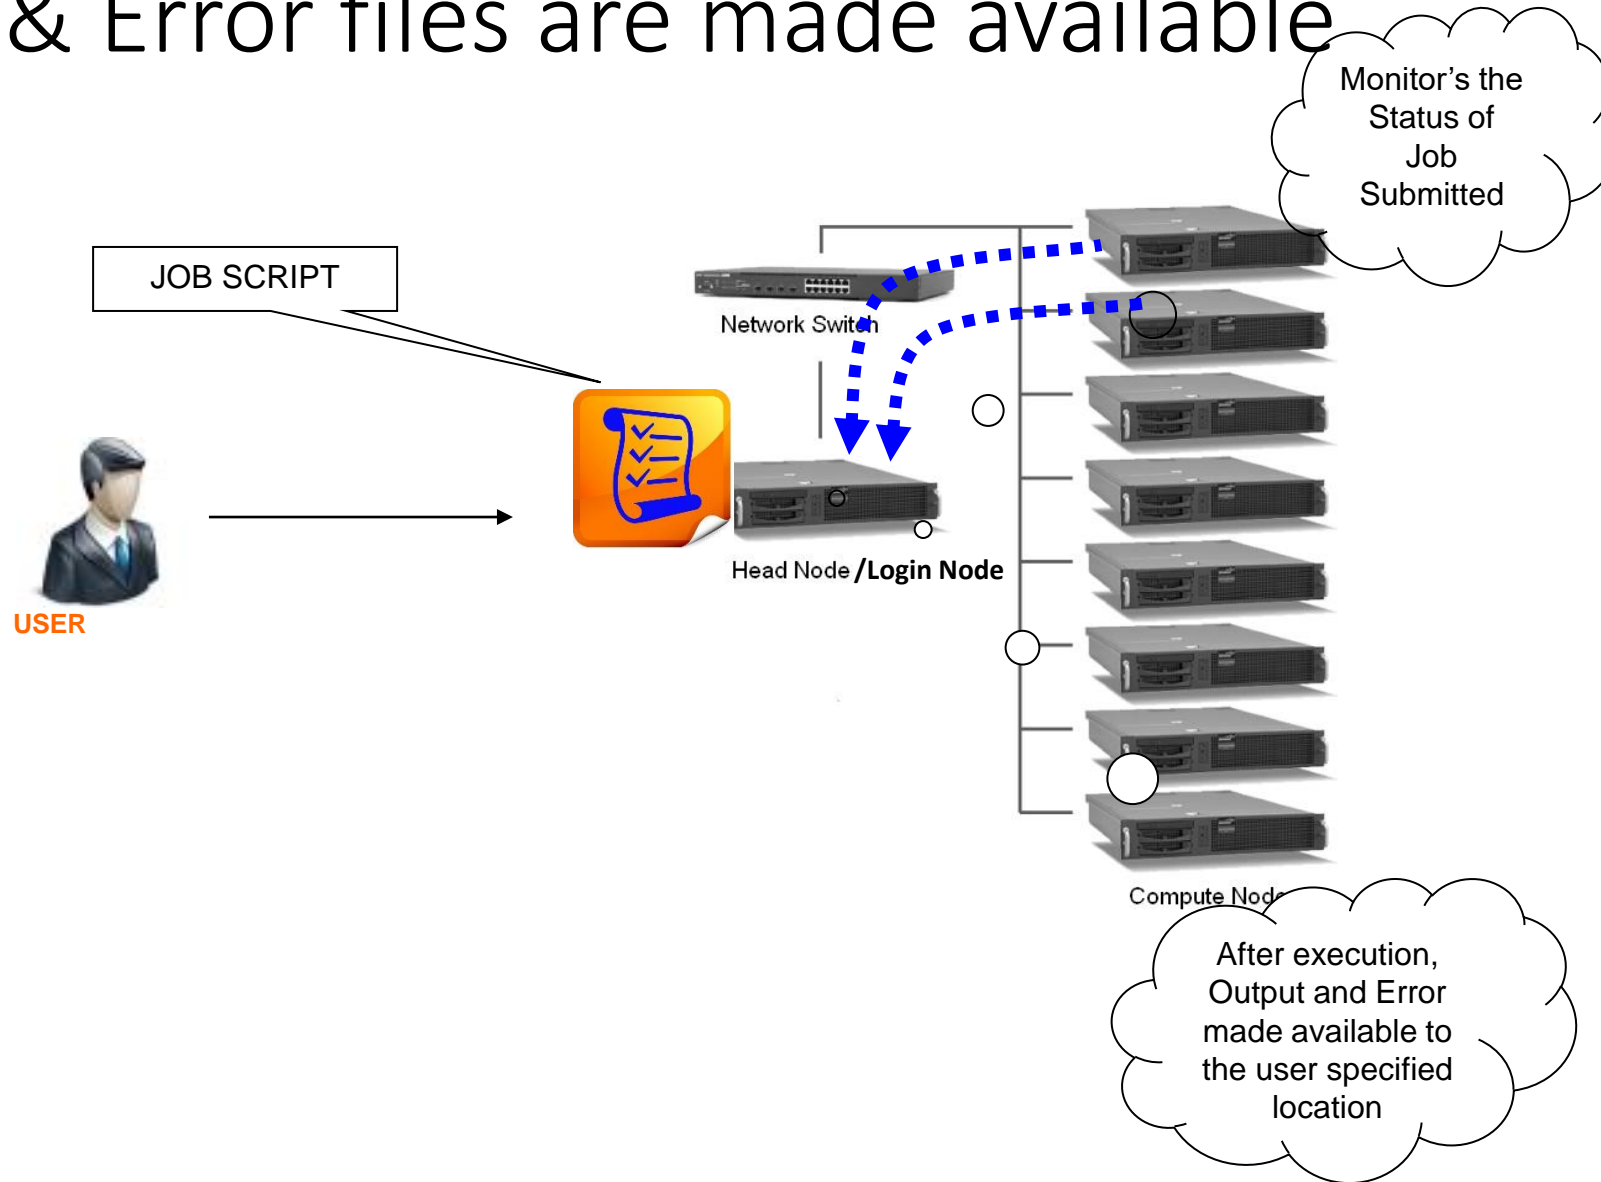
# Job Scheduling policies

- FCFS
  - Latest job submitted is added to the bottom of the queue.
  - No other job in the queue will run before the job that is at the top of the queue.
  - Top job waits in the queue until enough jobs finish to free up the resources that it needs.
- Multi-priority queues
  - Clusters with heterogeneous resources or job mixes configure multiple job queues for separation
  - Queues can be defined to support different job sizes, schedule to specific cluster resources
  - Queues can have different priorities to allow different categories of jobs to be scheduled differently
- Back-filling:
  - If the scheduler has the intelligence to launch jobs lower in the queue, on resources that are currently idle, it is called a back-fill scheduler.
  - The back-fill scheduler follows a strict rule to only schedule lower priority jobs on idle resources if it will not delay the start of the top priority job.
- Fair-share:
  - A method to allocate resource shares to a user or groups of users or a project
  - A fair method for ordering jobs based on their usage history or criteria based on pricing
  - The job to be run next is selected from the set of jobs belonging to the most deserving entity
- Preemptive:
  - A job with higher priority can signal currently running job to stop and release resources to allow the high priority job to run.
  - Necessary requirement for pre-emption is support for job checkpoints and restart ability.

# Job Execution in Compute Nodes



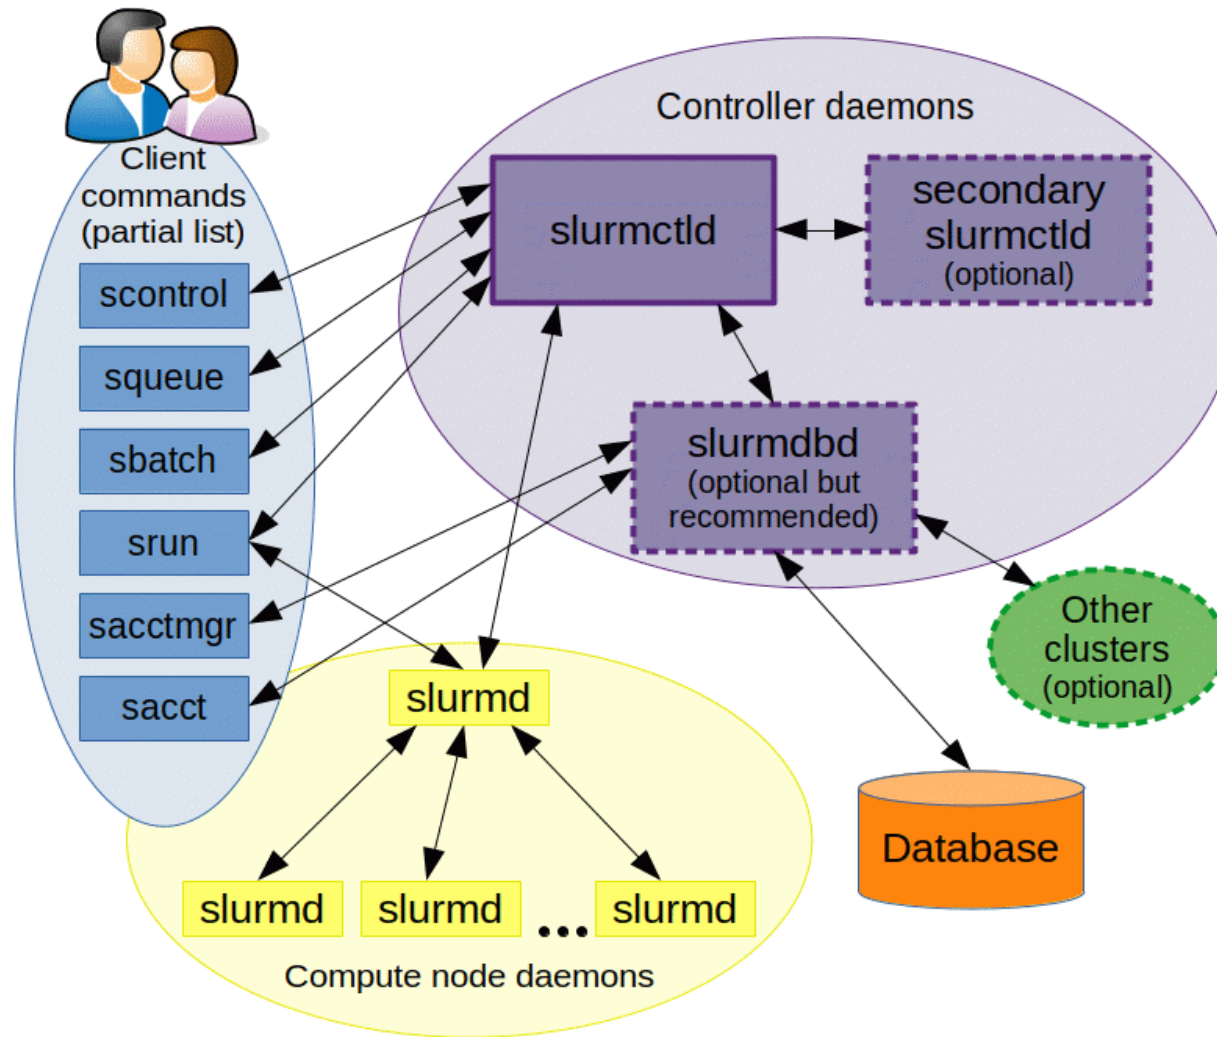
# Output & Error files are made available



# Local Resource Manager - SLURM

- When you login to HPC cluster, you land on Login Nodes
  - Login nodes are not meant to run jobs
  - These are used to submit jobs to Compute Nodes.
- To submit job on the cluster, you need to write a scheduler job script
- SLURM - Simple Linux Utility for Resource Management
- It is a local manager that provides a framework for job queues, allocation of compute nodes, and the start and execution of jobs.

# Local Resource Manager – SLURM - Components



# Local Resource Manager – SLURM - Commands

The list and descriptions of the mostly used Slurm commands

- **sbatch** <script> To submit the job on HPC cluster
- **squeue** To see the status of all jobs submitted on the cluster  
**squeue -u <user name>** To see status of user's jobs only. Also shows job-id.
- **sinfo** Provides the basic information about the resources on HPC cluster such as
  - Partitions/queue such as for cpu / gpu / high memory nodes
  - Number of nodes for each type and their numbering/names
  - State of the nodes

# Local Resource Manager – SLURM - Commands

- **scancel <job ID>** To delete the submitted jobs
- **scontrol show job <Job ID>** shows detailed information about a specific job or all jobs if no job id is given
- **sacct -u <username>** showing accounting stats after job ends
- **srun** To get resources in interactive mode
  - `srun --nodes=1 --ntasks-per-node=1 --time=00:05:00 --pty bash`

# SLURM: Sample Job Script for Serial Jobs

```
#!/bin/bash
#SBATCH -J TestJob
#SBATCH -N 1 // number of nodes
#SBATCH --ntasks-per-node=1 // number of cores per node
#SBATCH --output=3mm.out // name of output file
#SBATCH --error=3mm.err // name of error file
#SBATCH --partition=standard // partition or queue name
#SBATCH --time=01:00:00 // time required to execute the program
```



## SLURM: Sample Job Script for Parallel Jobs on CPU

```
#!/bin/bash
#SBATCH -N 1 // number of nodes
#SBATCH --ntasks-per-node=40 // number of cores per node
#SBATCH --output=3mm.out // name of output file
#SBATCH --error=3mm.err // name of error file
#SBATCH --partition=standard // partition or queue name
#SBATCH --time=01:00:00 // time required to execute the program

export OMP_NUM_THREADS=40
```

## SLURM: Job Script for Parallel Jobs on GPUs

```
#!/bin/bash
#SBATCH -N 1 // number of nodes
#SBATCH --ntasks-per-node=40 // number of cores per node
#SBATCH --output=3mm.out // name of output file
#SBATCH --error=3mm.err // name of error file
#SBATCH --time=01:00:00 // time required to execute the program
#SBATCH --gres=gpu:2 // request use of GPUs on compute nodes
#SBATCH --partition=gpu // partition or queue name

export OMP_NUM_THREADS=40
```

# PARAM Utkarsh

838TF



## CPU Only Compute Nodes

- ✦ 107 Nodes
- ✦ 5136 Cores
- ✦ Compute power of Rpeak 476.6 TFLOPS
- ✦ Each Node with
  - ✦ 2 X Intel Xeon Cascadelake 8268, 24 cores, 2.9 GHz, processors
  - ✦ 192 GB memory
  - ✦ 480 GB SSD

## GPU Compute Nodes

- ✦ 10 Nodes
- ✦ 400 CPU Cores
- ✦ 102400 CUDA Cores
- ✦ Rpeak CPU 32 TFLOPS + GPU 156 TF
- ✦ Each Node with
  - ✦ 2 X Intel Xeon Skylake 6248, 20 cores, 2.5 GHz, processors
  - ✦ 192 GB Memory
  - ✦ 2 x NVIDIA V100 SXM2 GPU Cards
  - ✦ 480 GB SSD


## High Memory Compute Nodes

- ✦ 39 Nodes
- ✦ 1872 Cores
- ✦ Compute power of Rpeak 173.7 TFLOPS
- ✦ Each Node with
  - ✦ 2 X Intel Xeon Cascadelake 8268, 24 cores, 2.9 GHz, processors
  - ✦ 768 GB Memory
  - ✦ 480 GB SSD

# System Details

| Parameter    | CPU only(75)           | GPU Nodes(10)   | GPU Ready(32)          | HM Nodes(39)           |
|--------------|------------------------|-----------------|------------------------|------------------------|
| Processor    | 2 x Xeon platinum 8268 | 2 x Xeon G-6248 | 2 x Xeon platinum 8268 | 2 x Xeon platinum 8268 |
| Cores        | 48                     | 40              | 48                     | 48                     |
| Speed        | 2.9 GHz                | 2.5 GHz         | 2.9 GHz                | 2.9 GHz                |
| Memory       | 192 GB                 | 192 GB          | 192 GB                 | <b>768 GB</b>          |
| HDD          | 480GB SSD              | 480GB SSD       | 480GB SSD              | 480GB SSD              |
| Total cores  | <b>3600</b>            | <b>400</b>      | <b>1536</b>            | <b>1872</b>            |
| Total Memory | 14400 GB               | 1920 GB         | 6144 GB                | 29952 GB               |
|              | -                      | 2 x NVIDIA V100 | -                      | -                      |

## PARAM Utkarsh Architecture Diagram



# Login - access methods

## System Access

### Accessing the cluster

- The cluster can be accessed through login nodes , which allows users to login.
- You may access login node through ssh.

## Remote Access

### a) Using SSH in Windows

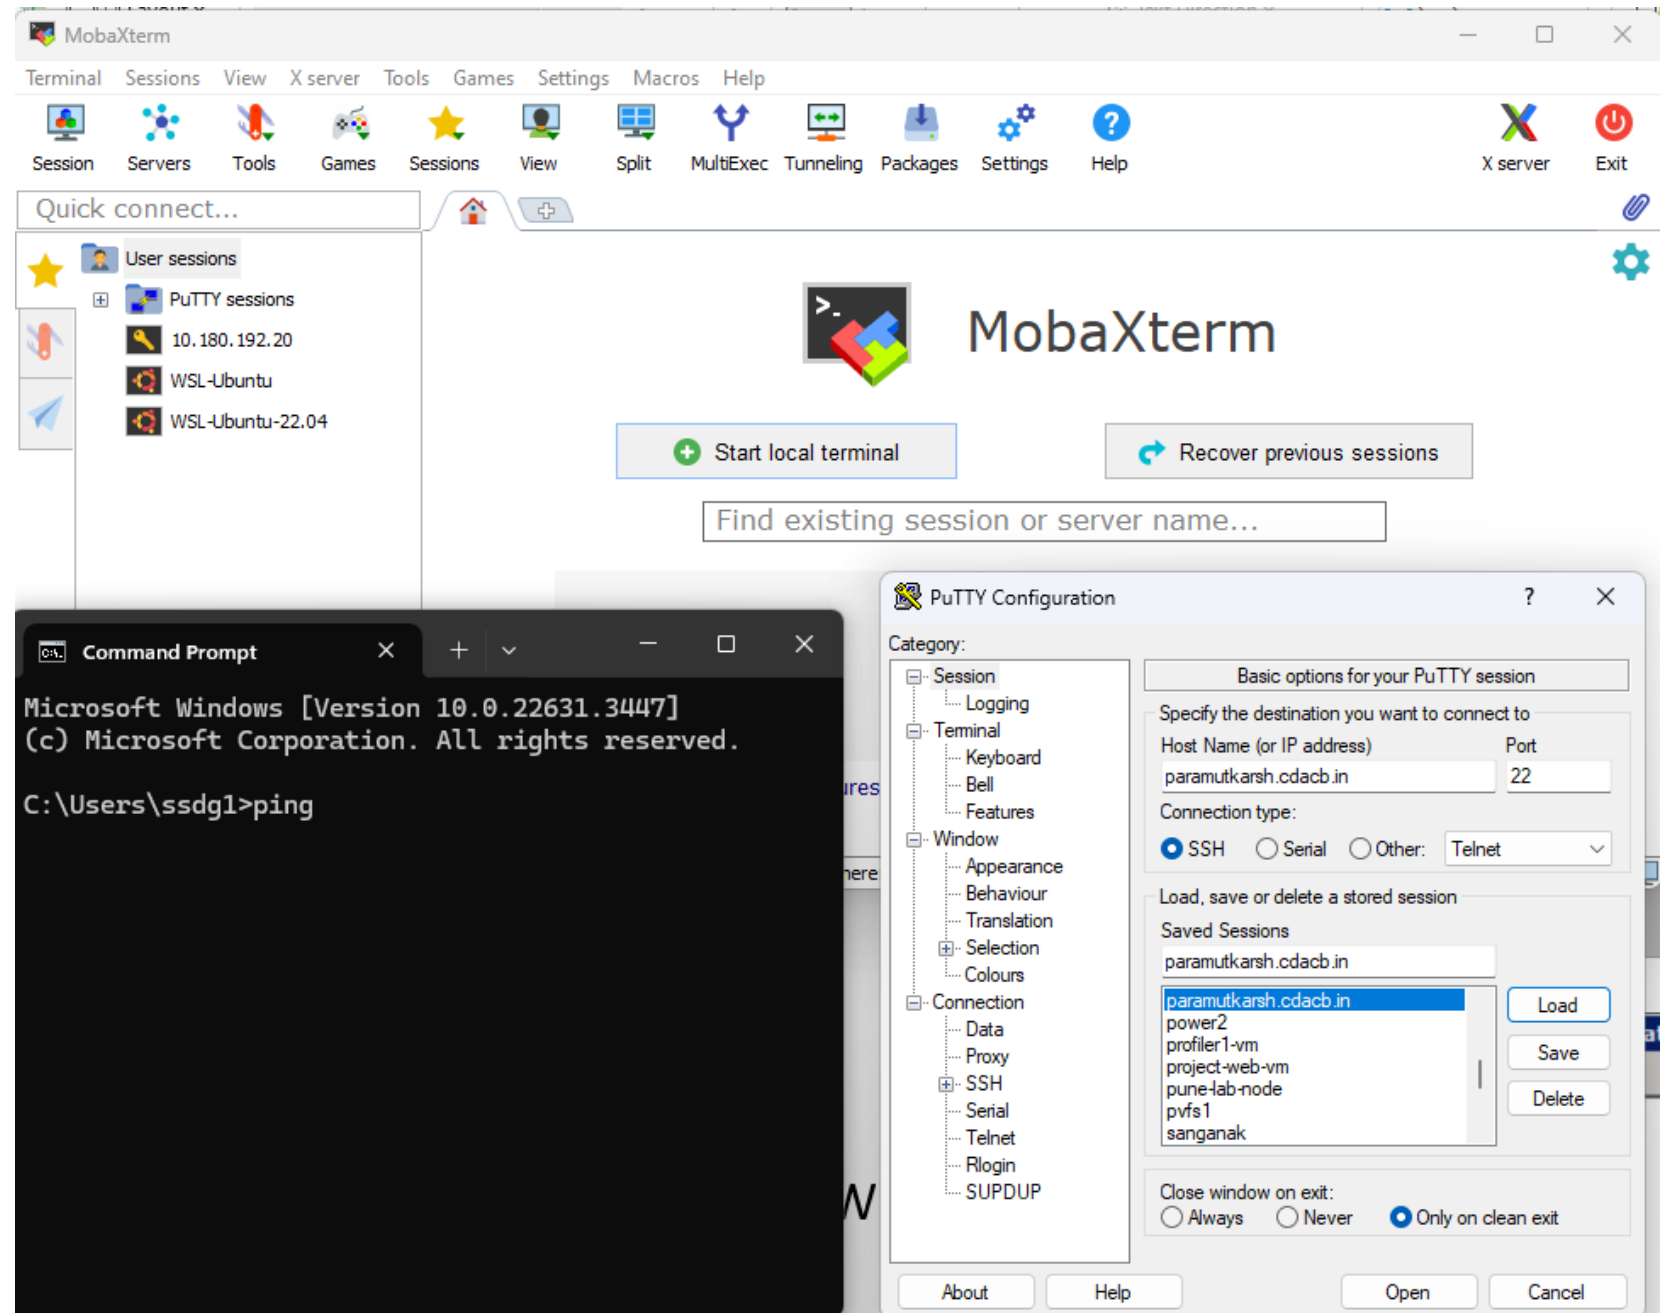
- To access PARAM Utkarsh you need to “ssh” the login server. PuTTY is the most popular open source “ssh” client application for Windows

### b) Using SSH in Mac or Linux

- Both Mac and Linux systems provide a built-in SSH client, so there is no need to install any additional package. Open the terminal, connect to an SSH server by typing the following command:

ssh [user1@paramutkarsh.cdacb.in](https://user1@paramutkarsh.cdacb.in)

# SSH Clients



# Working Environment

- In Linux, on the Unix core utilities are in your command-**PATH** by default.
- In Linux, only the default system libraries are in your **LD\_LIBRARY\_PATH**
- The **module** system allows users to easily update their working environment, to include specific codes, versions, compilers, and libraries.



# Module Utility

## How to set Environment ?

- By default no application is set in your environment. User must explicitly set required ones
- **module** is the utility (also command name) to enable use of applications/ libraries / compilers available on the HPC cluster.
- Module structure on Cluster
  - apps/<application name>/version :Applications available on the cluster
  - compiler/<compiler name>/version :Compilers available on the cluster
  - lib/<library name>/version :Available libraries
  - ....

# Module Utility

➤ Some Important commands:

- **module avail** To see the available software installed on HPC system
  - list of precompiled applications
  - different compilers and libraries (compilers include GNU, Intel, PGI)
- **module list** Shows the currently loaded modules in your shell
- **module load <Name of the module>**
  - *module load intel/2020*  
(to set Intel compilers version 2020 environment)
  - *module load namd/intel2018/2.12*  
(to set NAMD app version 2.12 environment)

# Module Utility

- **module unload <Name of the module>** This will remove all environment setting related to loaded previously
- **module purge** To clear all the loaded modules

## Transferring files between local machine and HPC cluster

- When a user wishes to transfer data from their local system (laptop/desktop) to HPC system They can just use “scp” command on their terminal
- The command shown below can be used for effecting file transfers  
`scp -r <path to the local data directory> <your username>@<IP of ParamUtkarsh>:  
<path to directory on HPC where to save the data>`

### Example:

***scp -r /dir/dir/file saurabh@<cluster IP/Name>:/home/Saurabh***

- Command could be used to transfer data from HPC system to your local system  
`scp -r <path to directory on HPC> <your username>@<IP of local system>:<path to the local data directory>`  
***scp -r /home/saurabh saurabh@<local system IP/Name>:/dir/dir/file***

# SLURM Job Arrays

- Job arrays allow you to leverage SLURM's ability to create multiple jobs from one script.
- For example, instead of having 5 submission scripts to run the same job with different arguments, you can have one script to run the 5 jobs at once.
- Many of the situations where this is useful include:
  - Running the same analysis program multiple times against different files or data sets.
  - Running the same program multiple times with different arguments.
  - Running a single program multiple times to analyzing a single data file.


# SLURM Job Arrays

- #SBATCH --array=*indexes*

|               |                             |
|---------------|-----------------------------|
| 1-10          | 1,2,3,4,5,6,7,8,9,10        |
| 2-20:2        | 2,4,6,8,10,12,14,16,18,20   |
| 1,3,5,7,11,21 | 1,3,5,7,11,21               |
| 2-20%2        | 2,4 then 6,8 then 10,12 ... |

# Dependent Jobs/Workflows

- The job dependency feature of SLURM is useful when you need to run multiple jobs in a particular order.
- A standard example of this is a workflow in which the output from one job is used as the input to the next



# Dependent Jobs Submission

- Job A

```
#!/bin/bash
#SBATCH --job-name=JobA
#SBATCH --time=00:05:00
#SBATCH --ntasks=1
#SBATCH --output=JobA.stdout
#SBATCH --error=JobA.stderr
echo "I'm job A"
echo "Sample job A output" > jobA.out
sleep 120
```



# Dependent Jobs Submission

- Job B

```
#!/bin/bash
#SBATCH --job-name=JobB
#SBATCH --time=00:05:00
#SBATCH --ntasks=1
#SBATCH --output=JobB.stdout
#SBATCH --error=JobB.stderr
echo "I'm job B"
echo "I'm using output from job A"
cat jobA.out >> jobB.out
echo "" >> jobB.out
echo "Sample job B output" >> jobB.out
sleep 120
```

# Dependent Jobs Submission

- Job C

```
#!/bin/bash
#SBATCH --job-name=JobC
#SBATCH --time=00:05:00
#SBATCH --ntasks=1
#SBATCH --output=JobC.stdout
#SBATCH --error=JobC.stderr
echo "I'm job C"
echo "I'm using output from job A"
cat jobA.out >> jobC.out
echo "" >> jobC.out
echo "Sample job C output" >> jobC.out
sleep 120
```

# Dependent Jobs Submission

- Job D

```
#!/bin/bash
#SBATCH --job-name=JobD
#SBATCH --time=00:05:00
#SBATCH --ntasks=1
#SBATCH --output=JobD.stdout
#SBATCH --error=JobD.stderr
echo "I'm job D"
echo "I'm using output from jobs B and C"
cat jobB.out >> jobD.out
echo "" >> jobD.out
cat jobC.out >> jobD.out
echo "" >> jobD.out
echo "Sample job D output" >> jobD.out
sleep 120
```

# Dependent Jobs Submission

- Submit Job A

```
[shamjith@login4 demo]$ sbatch JobA.slurm  
Submitted batch job 666898
```

- Submit Jobs B and C

```
[shamjith@login4 demo]$ sbatch -d afterok:666898 JobB.slurm  
Submitted batch job 666899
```

```
[shamjith@login4 demo]$ sbatch -d afterok:666898 JobC.slurm  
Submitted batch job 666900
```

- Submit Job D

```
[shamjith@login4 demo]$ sbatch -d afterok:666899:666900 JobD.slurm  
Submitted batch job 666901
```

# Job Script - Best Practices

- Keep unique copies of the stdout and stderr

```
#SBATCH -o jobname.%j.o  
#SBATCH -e jobname.%j.e
```

- echo commands back

```
#!/bin/bash -x  
set -x
```

- print statements

```
input=file1.inp  
echo $input
```

- print runtime environment

```
env
```

- make unique directories

```
mkdir -pv /home/$USER/${SLURM_JOB_ID}.${input}
```

# Types of Errors - Overview

- Scheduler (SLURM)
- Syntax
- Memory
- Storage
- File access
- Network
- Parallel communication

# Types of Errors - SLURM

- Scheduler (SLURM)
  - errors executing commands (sbatch, squeue)

*sbatch: error: Batch job submission failed: Unable to contact slurm controller*

*squeue: error: slurm\_receive\_msg: Socket timed out on send/rcv operation*  
*slurm\_load\_jobs error: Socket timed out on send/rcv operation*

slurmctld process may be getting high volume of requests

# Types of Errors - SLURM

- Scheduler (SLURM)
  - job states (ouput from squeue or sacct)

|   | State | Description |   |
|---|-------|-------------|---|
| 3 | CA    | CANCELLED   | Cancelled by the user or sysadm via scancel                 |
|   | CD    | COMPLETED   | All processes on all nodes completed with exit code = 0     |
|   | CG    | COMPLETING  | Job done, but processes on some nodes may still be active   |
|   | F     | FAILED      | Job terminated with non-zero exit code                      |
|   | NF    | NODE_FAIL   | Job terminated due to failure of one of the allocated nodes |
| 1 | PD    | PENDING     | Job is awaiting resource allocation                         |
|   | PR    | PREEMPTED   | Job terminated due to preemption                            |
| 2 | R     | RUNNING     | Job currently has an allocation                             |
|   | S     | SUSPENDED   | Execution has been suspended and CPUs                       |
|   | TO    | TIMEOUT     | Job terminated upon reaching its time limit (-t D-HH:MM)    |

**Normal state changes are:**



# Types of Errors - Syntax

- Syntax
  - job script

```
#!/bin/bash
#SBATCH -N 1
#SBATCH -n 1
#SBATCH -t 1:00:00
#SBATCH --mem=4000
#SBATCH -partition general
```

*sbatch: error: Invalid argument: general*

```
# This is a Job Script for Syntax Errors
input file1.txt

echo $input
```

*/var/slurmd/spool/slurmd/job70807187/slurm\_script: line 8: input: command not found*

# Types of Errors - Memory

- Memory
  - out of memory

*slurmstepd: error: Exceeded step memory limit at some point.*

- malloc failure
  - C function that allocates bytes of memory and returns a pointer to the allocated memory
- SIGSEGV, segfault or segmentation violation
  - arise primarily due to errors in use of pointers for virtual memory addressing, particularly illegal access.

# Types of Error - Storage

- Storage
  - out of space on device

*cp: closing `pgm\_simulation.out': No space left on device*

- out of space on filesystem quota
- out of inodes / file descriptors

*cp: cannot create regular file `simulation.sh': Disk quota exceeded*

# Types of Errors – File Access

```
# This is a Job Script for Syntax Errors
input=/n/home_rc/pedmon/a.out

cat $input
mpirun a.out
```

- File access
  - no permission to read/write

*/home/user/a.out: Permission denied.*

- file or library not found

*/home/user/a.out: error while loading shared libraries: libm1.so.0: cannot open shared object file: No such file or directory*

- command not found

*/var/slurmd/spool/slurmd/job2255/slurm\_script: line 16: mpirun: command not found*

# Performance Analysis

# Speedup

- Speedup measures change in running time due to parallelism. The number of PEs is denoted here by  $n$ .
- Based on running times,  $S(n) = t_s/t_p$ , where
  - $t_s$  is the execution time on a single processor, using the fastest known sequential algorithm
  - $t_p$  is the execution time using a parallel processor.

$$\text{Speedup} = \frac{\text{Sequential execution time}}{\text{Parallel execution time}}$$

# Amdahl's Law

This law, formulated by computer architect Gene Amdahl, quantifies the maximum potential speedup achieved by parallelizing a computation while considering the fixed portion of the task that cannot be parallelized.

Let  $f$  be the fraction of operations in a computation that must be performed sequentially, where  $0 \leq f \leq 1$ . The maximum speedup  $\psi$  achievable by a parallel computer with  $n$  processors is

$$\psi \equiv S(n) \leq \frac{1}{f + (1-f)/n} \leq \frac{1}{f}$$

# Example 1

- 95% of a program's execution time occurs inside a loop that can be executed in parallel. What is the maximum speedup we should expect from a parallel version of the program executing on 8 CPUs?

$$\psi \leq \frac{1}{0.05 + (1 - 0.05) / 8} \cong 5.9$$



# Gustafson's Law

- Gustafson's Law, formulated by computer scientist John L. Gustafson, provides an alternative perspective on the potential performance improvements achievable through parallel computing.
- Gustafson's Law considers the workload's scalability and disregards the idea of having a fixed portion of the workload.
- It states that as the computing resources (typically the number of processors) increase, the problem size and workload can be scaled up to utilize the available resources effectively.
- This shows that even more significant problems can be solved with parallel computing, which is more efficient as it focuses on keeping the problem size proportional to the available resources.
- $$S(n) = n + (1 - f)n$$

$S(n)$  is the theoretical speedup with  $n$  processors, and  $f$  is the serial fraction of the job

# Granularity

- For parallel tasks we talk about the *granularity* – size of the computation between *synchronization points*
  - Coarse – heavyweight processes + IPC (interprocess communication, like MPI)
  - Fine – Instruction level (eg. SIMD)
  - Medium – Threads + [message passing + shared memory ]
- Computation to Communication Ratio
  - $(\text{Computation time}) / (\text{Communication time})$
  - *Increasing* this ratio is often a key to good efficiency

# Communication Overhead

- Another important metric is *communication*
  - Measure time spent on communication that *cannot* be spent on computation
- *Overlapped Messages* –
  - portion of message lifetime that can occur concurrently with computation

Debugging

# Parallel vs. serial

- Parallel programming is more difficult than serial programming each step of the way:
  - Designing stage
  - Coding
  - Debugging
  - Profiling
  - Maintenance

# Parallel bugs

- In addition to usual, “serial” bugs, parallel programs can have “parallel-only” bugs, such as
  - Race conditions
    - when two or more threads access shared data concurrently, and the final outcome depends on the order of execution.
  - Deadlocks
    - when two or more tasks or threads are blocked, each waiting for the other to release a resource or respond to a signal.

# Race condition

- In OpenMP, race conditions result from misuse of shared variables, when
  - A variable is mistakenly labeled as shared (where in fact it needs to be private), or
  - A variable is correctly labeled as shared, but the access to the variable wasn't properly protected (serialized)
  - For example, if two threads execute the code concurrently, one thread might read the value of  $b$  while another thread is updating it, leading to inconsistent results
    - $a = b + c;$
    - $b = a + c;$
    - $c = b + a;$

# Race condition

```
#include <omp.h>
#include <stdio.h>
int main(int argc, char *argv[]) {
    int numt, tid;
    #pragma omp parallel
    {
        // All threads are accessing and modifying the shared variables numt and tid
        numt = omp_get_num_threads();
        tid = omp_get_thread_num();
        printf("Hello world from thread %d of %d\n", tid, numt);
    }
    return 0;
}
```

Hello world from thread 1 of 4  
Hello world from thread 0 of 4  
Hello world from thread 1 of 4  
Hello world from thread 2 of 4



# Race condition-Solution

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[]) {
    int numt, tid;

    #pragma omp parallel private(tid) shared(numt)
    {
        numt = omp_get_num_threads();
        tid = omp_get_thread_num();

        printf("Hello world from thread %d of %d\n", tid, numt);
    }

    return 0;
}
```

Hello world from thread 1  
of 4 Hello world from  
thread 0 of 4 Hello world  
from thread 3 of 4 Hello  
world from thread 2 of 4

# Example of race condition

Shared variable incremented by multiple threads without protection.

...

```
int main() {  
    int shared_var = 0;  
    #pragma omp parallel num_threads(4)  
    {  
        #pragma omp for  
        for (int i = 0; i < 1000; i++) {  
            shared_var++;  
        }  
    }  
    printf("Shared variable value: %d\n", shared_var);  
    return 0;  
}
```

Use synchronization constructs (e.g., critical sections, locks, atomic operations) to protect access to shared variables

...

```
int main() {  
    int shared_var = 0;  
    #pragma omp parallel num_threads(4)  
    {  
        #pragma omp for  
        for (int i = 0; i < 1000; i++) {  
            #pragma omp critical  
            {  
                shared_var++;  
            }  
        }  
    }  
}
```

# Deadlocks

- It happens when thread(s) lock up while waiting on a locked resource that will never become available
- The sign of a deadlock: the program hangs (always or sometimes) when reaching a certain point in the code

# Deadlocks (cont.)

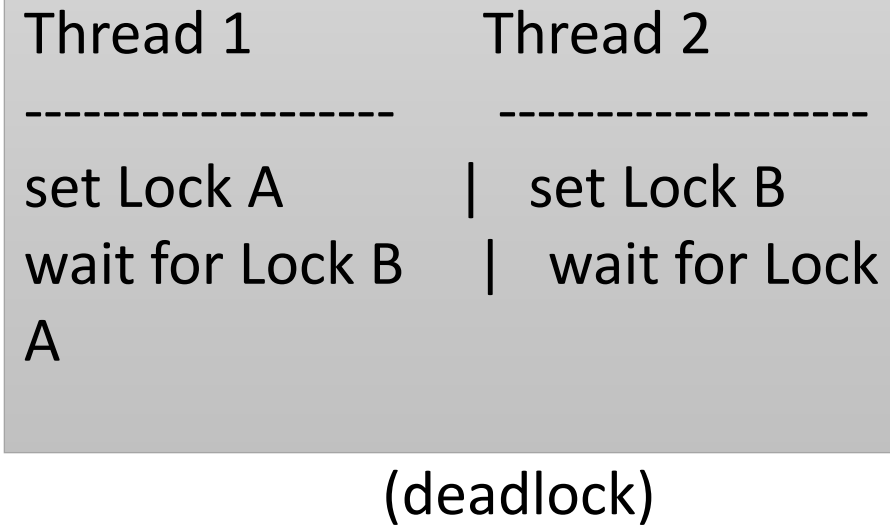
## ➤ Prevention strategies:

- Be very careful with conditional clauses using `threadID` as an argument, as common OpenMP constructs (`for/do`, `single`) require all the threads in the team reaching them.
- Communications between threads (using a shared variable) have to use “flush” pragma, on both writing and reading sides. – Don't forget to unset locks after setting them.

# Example of a deadlock

we can ensure that the locks are always acquired and released in the same order in all sections.

```
if (thread_id == 0) {  
    omp_set_lock(&locka);  
    printf("Thread 0 locked locka\n");  
    omp_set_lock(&lockb);  
    printf("Thread 0 locked lockb\n");  
    omp_unset_lock(&lock1);  
    omp_unset_lock(&lock2);  
} else {  
    omp_set_lock(&lockb);  
    printf("Thread 1 locked lockb\n");  
    omp_set_lock(&locka);  
    printf("Thread 1 locked locka\n");  
    omp_unset_lock(&lock1);  
    omp_unset_lock(&lock2);  
}
```



# Overview of GDB

- “GNU Debugger”
- A debugger for several languages, including C and C++
- It allows you to inspect what the program is doing at a certain point during execution.
- Errors like segmentation faults may be easier to find with the help of gdb
- [http://sourceware.org/gdb/current/onlinedocs/gdb toc.html](http://sourceware.org/gdb/current/onlinedocs/gdb%20toc.html) - online manual

# Preparing your program

- Additional step when compiling program
- Add a `-g` option to enable built-in debugging support (which gdb needs)
- `-g` which tells the compiler to generate symbolic information required by any debugger.
- For examples
- `gcc -g -o0 -fopenmp -o code code.c`

# GDB Commands

- Related Gdb Commands:
  - List: list the source code and each execution's corresponding line number
  - Break linenum: set breakpoint at the linenum
  - Run argv: run the execution code with the parameter argv
  - Next: execute the next line of code
  - Backtrace: show trace of all function calls in stack
  - **Info frame**: List address, language, address of arguments/local variables and which registers were saved in frame.
    - This will show where the return address is saved
    - Return address is in Register EIP
    - Calling stack pointer is in Register EBP
  - x &variable: show the address and value of a local variable (in hex format)
  - **x address**: print binary representation of bytes of memory pointed to by address.



# Example of Using GDB - Serial

```
#include <stdio.h>

void foo(char * input){
    int a1=11;
    int a2=22;
    char buf[7];
    strcpy(buf, input);
}

void main(int argc, char **argv){
    foo(argv[1]);
}
```

# Starting up gdb

- Just try “gdb” or “gdb program.x” to access the GDB prompt
  - *(gdb)*
- To execute the program
  - *(gdb) run*
- If the program encounters issues, GDB provides valuable information.
- For example: Segmentation Fault

```
Program received signal SIGSEGV, Segmentation fault.  
0x00000000004008bd in fib (i=5) at fibb.c:21  
21          *ptr = 42; // Dereferencing a null pointer
```

# Setting breakpoints

- Breakpoints can be used to stop the program run in the middle, at a designated point.
- Useful for investigating issues or examining specific section of code
- Setting Breakpoints - File and Line:
  - Syntax: (gdb) break file.c:6
  - Set a breakpoint at line 6 in the file "file.c."
- Setting Breakpoints - Function:
  - Syntax: (gdb) break my\_func
  - Set a breakpoint at the beginning of the function "my\_func."

# Now what ?

- After setting a breakpoint, use the *run* command to execute the program until it reaches the specified breakpoint.
- Stepping to the Next line
  - \$(gdb) next
- Stepping into functions
  - \$(gdb) step
- Listing Source code
  - \$(gdb) list

# Other useful commands

- Inspect variable values during the paused state.
  - `$(gdb) print <variable>`
- Continue execution after inspection.
  - `$(gdb) continue`
- Examining Stack Status:
  - `$(gdb) backtrace`

# OpenMP Debugging with GDB

- Introduction to GDB as a Debugger for OpenMP Programs:
  - GDB (GNU Debugger) is a powerful command-line debugger for C/C++ programs on UNIX systems.
  - Widely used for debugging OpenMP programs to identify and resolve issues.
- Commands and Techniques for Debugging OpenMP Code with GDB:
  - Start up GDB to debug a program
  - `$ OMP_NUM_THREADS=2 gdb ./application.exe`

# OpenMP Debugging with GDB

## Thread-specific Commands in GDB

- (gdb) info thread- Prints out information about all current threads.
  - Helpful for understanding the status of threads during program execution, especially in parallel sections.
- (gdb) thread -Prints the current thread
- (gdb) thread <thread\_no> - switches to specific thread.
  - Allows the user to switch between threads and focus debugging efforts on specific threads of interest.

# OpenMP Debugging with GDB

- thread apply- Sending GDB Commands to Multiple Threads
  - Useful for applying debugger commands uniformly across multiple threads.
  - (gdb) thread apply 2-4 continue
  - (gdb) thread apply all print tid
- Scheduling Locking:
  - (gdb) set scheduler-locking on
  - (gdb) set scheduler-locking off
  - (gdb) show scheduler-locking



# Debuggers

- GDB
- LLDB
- TotalView
- DDT

# Profiling

# What is Profiling


- Application profiling is the process of analysing and measuring the performance of a software application in order to identify and diagnose performance issues or bottlenecks.
- Profiling tools help developers to identify areas of code that are taking a long time to execute, or that are using an excessive amount of system resources like memory, CPU or disk I/O.
- The purpose of profiling is to identify any bottlenecks or areas of the application that can be optimized to improve its performance, scalability, and reliability.
- By profiling their applications, developers can ensure that they are delivering high-quality software that performs well and meets the needs of their users.
- Profiling involves collecting data on various aspects of the application and its analysis.
- Implemented through either
  - Sampling based profiling
  - Instrumentation based Profiling

# Sampling Based Profiling

- With sampling, the executable is stopped at regular intervals. Each time it is halted, key information is gathered and stored.
- At specified intervals, the Sampling method collects information about the functions that are executing in your application. While the program is interrupted the profiler grabs a snapshot of its current state.
- Sampling is a low cost to the profiler and has little effect on the execution of the application being profiled.
- If you need accurate measurements of call times or are looking for performance issues in an application, then sampling based profiler are useful.
- Sampling has less accuracy in the number of calls

# Inclusive Vs Exclusive CPU Time

- This is an important concept in profiling tools
- The inclusive metric includes all callees underneath the caller - For example, all the CPU time accumulated when executing a function
- The exclusive metric excludes everything outside the caller - For example, the CPU time accumulated outside of calling other functions



| Function | Inclusive time | Exclusive time |
|----------|----------------|----------------|
| A        | 75             | 10             |
| B        | 20             | 20             |
| C        | 30             | 5              |
| D        | 15             | 15             |
| E        | 25             | 25             |

```
Void Alpha()  
{
```



30 samples

```
Beta();
```

```
}
```

```
Void Beta()  
{
```



50 samples

```
}
```

| Functions | Inclusive | Exclusive |
|-----------|-----------|-----------|
| Alpha     | 80        | 30        |
| Beta      | 50        | 50        |

# Instrumentation based Profiling

- The first and earliest type are instrumenting profilers.
- Data collection is done by tools that either injecting code into a binary file that captures timing information or by using callback hooks to collect and emit exact timing and call count information while an application runs.
- The instrumentation method has a high overhead when compared to sampling-based approaches.
- Instrumentation profiling is that you can get exact call counts on how many times your functions were called.
- Instrumentation can be achieved during
  - Compile time Instrumentation
  - Binary Instrumentation

# Performance Profilers

- TAU
- Vtune
- GProfNG
- HPCToolkit
- Likwid



# Petascale and Exascale Computing

# AIRAWAT - PARAM Siddhi AI



**AIRAWAT - PARAM Siddhi - AI of 8.5 Petaflops is the fastest Supercomputer in India and ranked at No. 90 position in 'TOP500 Supercomputer List – November 2023'**

# List of systems commissioned in Phase-I and Phase-II (24PF)

| S.No. | Institute Name                     | HPC System Name | Computing Power  | Year of Commissioning |
|-------|------------------------------------|-----------------|------------------|-----------------------|
| 1.    | IIT(BHU), Varanasi                 | PARAM Shivay    | 838TF            | 2019                  |
| 2.    | IISER, Pune                        | PARAM Brahma    | 1.7PF            | 2020                  |
| 3.    | IIT, Kharagpur                     | PARAM Shakti    | 1.66PF           | 2020                  |
| 4.    | JNCASR, Bangalore                  | PARAM Yukti     | 1.8PF            | 2020                  |
| 5.    | IIT, Kanpur                        | PARAM Sanganak  | 1.66PF           | 2020                  |
| 6.    | C-DAC, Pune (National AI Facility) | PARAM Siddhi    | 5.2PF/210PF (AI) | 2020                  |
| 7.    | IIT, Hyderabad                     | PARAM Seva      | 838TF            | 2021                  |
| 8.    | NABI, Mohali                       | PARAM Smriti    | 838TF            | 2021                  |
| 9.    | IISc, Bangalore                    | PARAM Pravega   | 3.3PF            | 2022                  |
| 10.   | C-DAC, Bangalore (MSME)            | PARAM Utkarsh   | 838TF            | 2021                  |
| 11.   | IIT, Roorkee                       | PARAM Ganga     | 1.66PF           | 2022                  |
| 12.   | IIT, Gandhinagar                   | PARAM Ananta    | 838TF            | 2022                  |
| 13.   | NIT, Trichy                        | PARAM Porul     | 838TF            | 2022                  |
| 14.   | IIT, Guwahati                      | PARAM Kamrupa   | 838TF            | 2022                  |
| 15.   | IIT, Mandi                         | PARAM Himalaya  | 838TF            | 2022                  |

# NSM Commissioned Sites

PARAM Shivay



PARAM Shakti



PARAM Brahma



PARAM Yukti



PARAM Sanganak




PARAM Pravega



PARAM Seva



PARAM Smriti



PARAM Himalaya



PARAM Kamrupa



PARAM Utkarsh



PARAM Ganga



PARAM Ananta



PARAM Porul



# Top 500 Supercomputers

| Rank | System  | Cores     | Rmax<br>(PFlop/s) | Rpeak<br>(PFlop/s) | Power<br>(kW) |
|------|---|-----------|-------------------|--------------------|---------------|
| 1    | <b>Frontier</b> - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE<br>DOE/SC/Oak Ridge National Laboratory<br>United States                  | 8,699,904 | 1,194.00          | 1,679.82           | 22,703        |
| 2    | <b>Aurora</b> - HPE Cray EX - Intel Exascale Compute Blade, Xeon CPU Max 9470 52C 2.4GHz, Intel Data Center GPU Max, Slingshot-11, Intel<br>DOE/SC/Argonne National Laboratory<br>United States | 4,742,808 | 585.34            | 1,059.33           | 24,687        |
| 3    | <b>Eagle</b> - Microsoft NDv5, Xeon Platinum 8480C 48C 2GHz, NVIDIA H100, NVIDIA Infiniband NDR, Microsoft<br>Microsoft Azure<br>United States  | 1,123,200 | 561.20            | 846.84             |               |
| 4    | <b>Supercomputer Fugaku</b> - Supercomputer Fugaku, A64FX 48C 2.2GHz, Tofu interconnect D, Fujitsu<br>RIKEN Center for Computational Science<br>Japan   | 7,630,848 | 442.01            | 537.21             | 29,899        |



# Frontier Exascale Supercomputer

1.1EF

700  
PETABYTES



40  
MEGAWATTS

8699904  
Cores

# Challenges

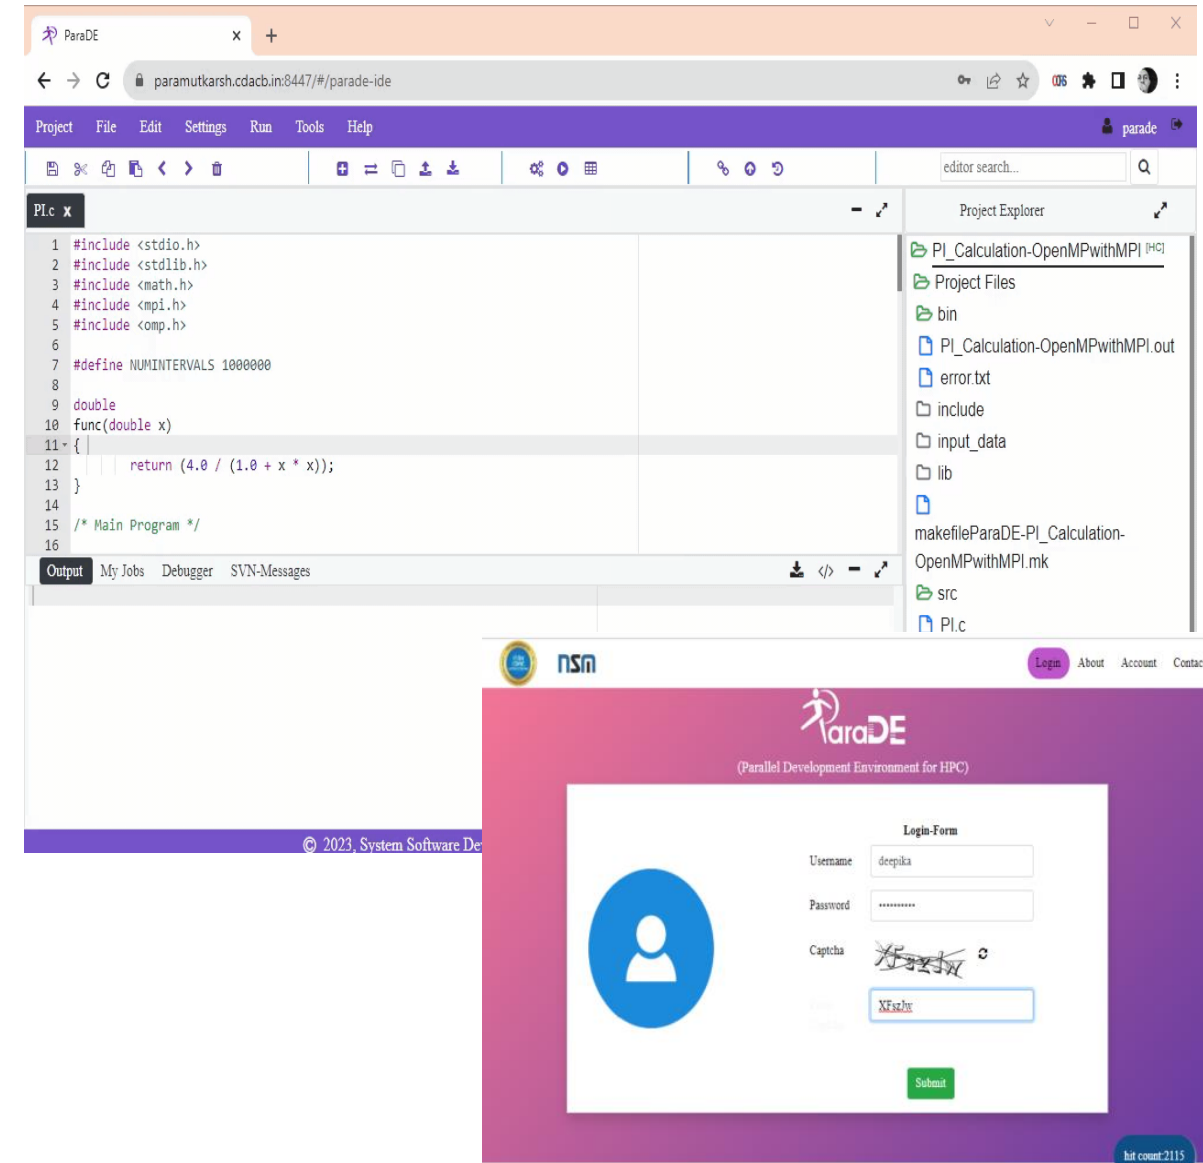
- Power consumption. ...
- Scalability. ...
- Heterogeneity
- Resilience. ...
- Programming methodologies and applications

C-DAC's System Software for HPC



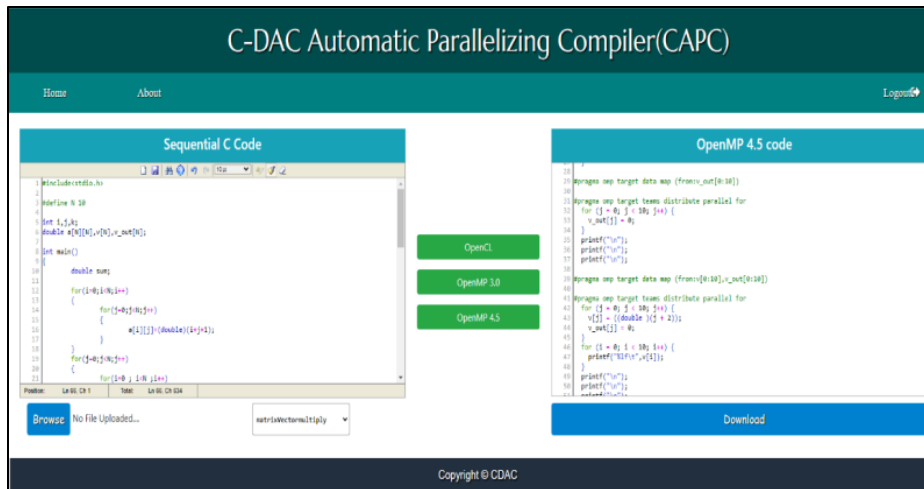

# ParaDE (Parallel Development Environment)

- It is an **IDE for HPC** which facilitates application development on HPC platform from anywhere and anytime
- Major Features
  - ✓ Support **multiple paradigms** like MPI, OpenMP, CUDA, OpenACC under a single project
  - ✓ **Automatic compilation**
  - ✓ **Assisted job submission**
  - ✓ Provision to **add custom** compilers, libraries and tools
  - ✓ **Version Control System**
  - ✓ Access to development tools (**debuggers, profilers, converters**) available on the platform



# CAPC - CDAC Automatic Parallelizing Compiler

- Automatically converts sequential programs to equivalent parallel programs for the target parallel architectures



## Features

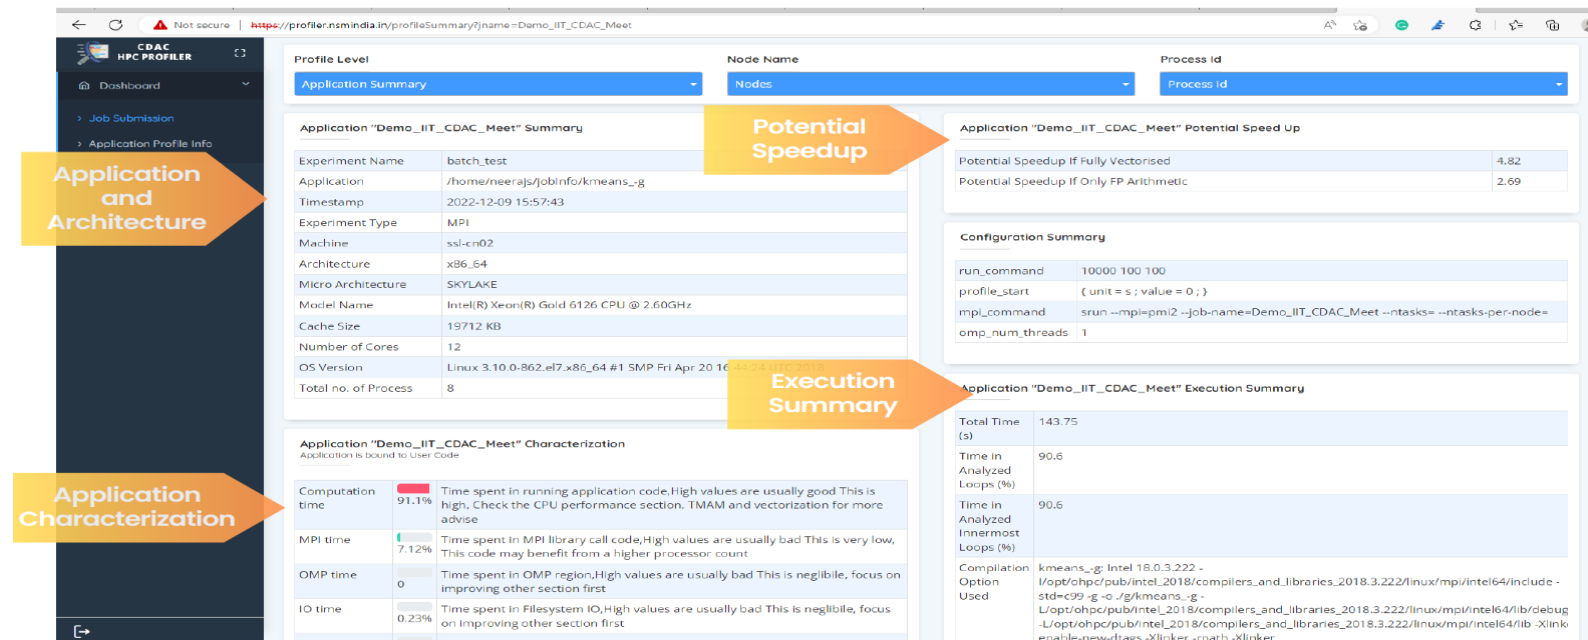
- Automatic parallelizes the code without user hints/inputs
- Support for multiple Parallel paradigms
- Human readable output
- Profitability estimate before parallelizing

# CHAP - CDAC HPC Application Profiler

- It provide insights on the application performance behaviour at the cluster level and process level

- Major Features

- ✓ Application **Performance Summary**
- ✓ **Guided** profiling
- ✓ **Hotspot** identification
- ✓ Multi-dimension analysis
- ✓ Minimal overhead in profiling
- ✓ Potential **Performance suggestions**
- ✓ **Integration with LRMs**
- ✓ Access via **web interface**

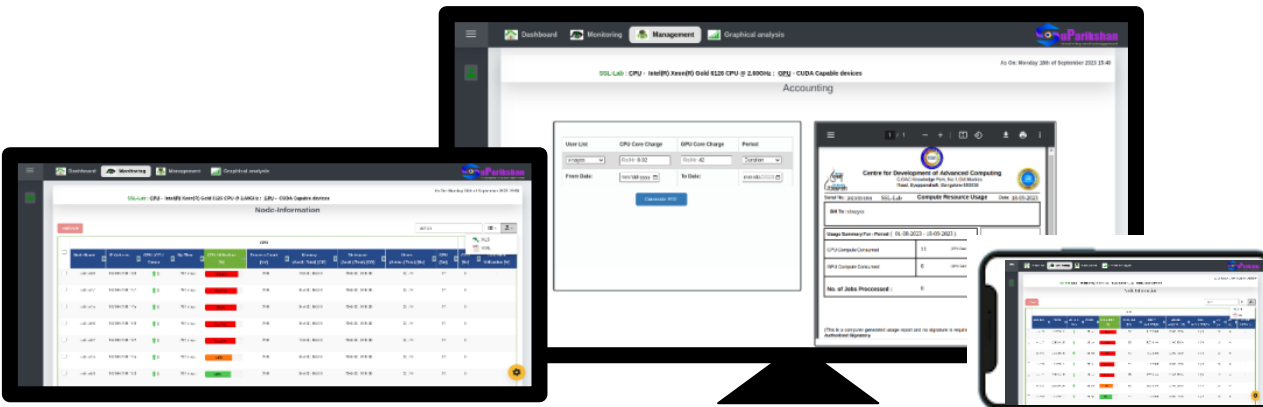
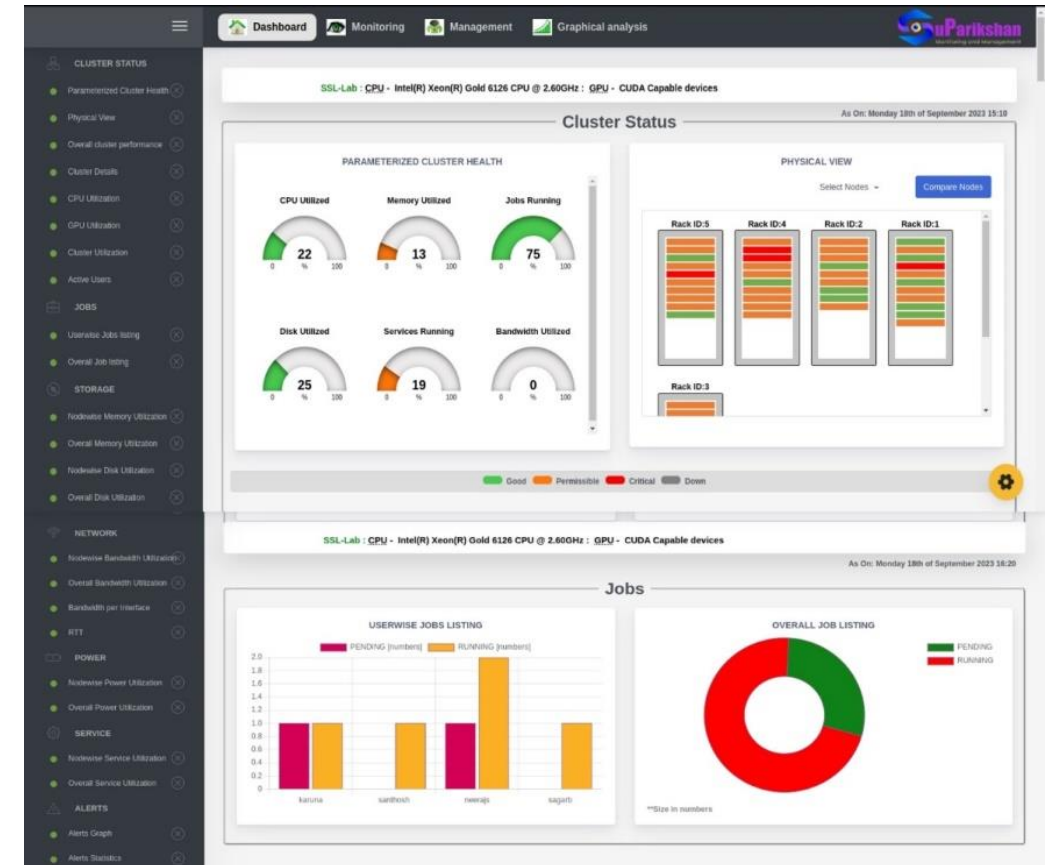


# SUPARIKSHAN

❑ Monitoring & management solution for heterogeneous HPC cluster. It is a pluggable, customizable and integrate with other third party tools


## ❑ Major Features

- ✓ Assist Administrators to manage heterogeneous infrastructure like CPU, GPU, FPGA etc
- ✓ Low Footprint software with responsive design
- ✓ Automatic error identification & ticketing system
- ✓ Supports monitoring customized scripts of admin
- ✓ Accounting of cluster resources and users



# HPC Dashboard

HPC Dashboard hides the complexities of high performance computing (HPC). It has a rich, easy to use interface for end-users and administrators, designed to increase productivity through its visual web-interface, powerful job submission and management feature, remote visualization, accounting & reporting, secure web console, file manager, administration panel and other workload functions.



## Features of HPC Dashboard

**Personalized  
Dashboard**

**Secured Web  
Console**

**Mad Libs  
Job submission**

**Accounting &  
Reporting**

**Advanced  
File manager**

**Interactive  
Voice**

**Integrated  
Help Desk**

**Single  
Sign-On**

**Resource  
Administration**

**Remote  
Visualization**

Questions